

Ch4. SQL

A. Basic Query

1. SELECT DISTINCT eliminate equivalent rows
2. FROM
3. WHERE AS / = LIKE string matching 'B-%B' - any one character, % 0 or more

B. Set Operations

4. UNION
5. INTERSECT
6. EXCEPT

C. Comparison Operators

1. IN
2. EXISTS
 $\sigma_{eid}(\text{employees}) - \sigma_{eid}(\text{dependent})$
3. op ANY
 $\sigma_{name}(\text{employees} \bowtie \text{participate} \bowtie \text{project})$
4. op ALL

Ch 5. Schema Refinement & Normal Forms

A. Motivation

<u>Id</u>	name	age	rating	Hourly_wages	Hours_worked
123-22-3666	Peter	48	8	10	40
231-31-5368	Paul	22	8	10	30
131-24-3650	Mary	35	5	7	30
434-26-3751	David	35	5	7	32
612-67-4134	Ada	35	8	10	40

Rating = 8 \rightarrow Hourly_wages = 10

Rating = 5 \rightarrow Hourly_wages = 7

★ Redundant Storage

~ the associations are repeated (storage is not used efficiently)

1. Update Anomalies

~ do not make a similar change to other tuples if a tuple is updated (may cause inconsistency)

2. Insertion Anomalies

~ not allow to insert a tuple unless we know that hourly wage and rating (otherwise, null values may exist)

3. Deletion Anomalies

~ if we delete all tuples with a given rating value (David and Mary), we may lose the associations

• Decomposition

~ replacing the relation schema by two or more relation schemas that each contains a subset of attributes

<u>Id</u>	name	age	rating	Hourly_wages	Hours_worked
123-22-3666	Peter	48	8	10	40
231-31-5368	Paul	22	8	10	30
131-24-3650	Mary	35	5	7	30
434-26-3751	David	35	5	7	32
612-67-4134	Ada	35	8	10	40

subset of attributes

<u>Id</u>	name	age	rating	Hours_worked
123-22-3666	Peter	48	8	40
231-31-5368	Paul	22	8	30
131-24-3650	Mary	35	5	30
434-26-3751	David	35	5	32
612-67-4134	Ada	35	8	40

rating	Hourly_wages
8	10
5	7

Functional dependency:
- rating determines Hourly_wages

B. Functional Dependencies FD

~ a kind of IC that generalizes the concept of a key

Let R be a relation schema and let X and Y be nonempty sets of attributes in R.

We say that an instance r of R satisfies the FD $X \rightarrow Y$
If the following holds for every pair of tuples t_1 and t_2 in r

If $t_1.X = t_2.X$, then $t_1.Y = t_2.Y$

The notation $t_1.X$ refers to the projection of tuple t_1 onto the attributes in X

α determines β

	X	Y
	rating	Hourly_wages
t1	8	10
t2	8	10
	5	7
	5	7
	8	10

If $t_1.X = t_2.X = 8$

then $t_1.Y = t_2.Y = 10$

student_ID	student_name	course_ID	course_name
111	Chan Tai Man	3170	Database
222	Wong Siu Ling	3170	Database
333	Tam Wai Ming	3160	Algorithms
111	Chan Tai Man	3160	Algorithms

- Trivially Preserved

~ $\alpha \rightarrow \beta$ is trivially preserved if any two rows never agree on α value

$course_ID \rightarrow course_name$ (NOT trivially preserved)

$student_ID, course_ID \rightarrow course_name$ (trivially preserved)

\therefore cannot find two rows with same $student_ID$ and $course_ID$

- Trivial Dependency

~ $\alpha \rightarrow \beta$ is a trivial dependency if $\beta \subseteq \alpha$

$student_name \rightarrow student_name$ (trivial dependency)

$student_name, course_name \rightarrow student_name$

- Keys

1. α is superkey for R iff $\alpha \rightarrow R$

2. α is candidate key for R iff $\alpha \rightarrow R$ and no γ that is a proper subset of α , $\gamma \rightarrow R$

- Reasoning about FDs

F - a set of functional dependencies

f - an individual functional dependency

f is implied by F if all functional dependencies in F is true, then f is true

e.g. $F = \{id \rightarrow did, did \rightarrow office\}$ implies $id \rightarrow office$

- Closure of F F^+

~ if the set of all FDs implied by F , then F is a closure of F , denoted as F^+

$$F = \{A \rightarrow B, B \rightarrow C, C \rightarrow D\}$$

$$F^+ = \{A \rightarrow B, B \rightarrow C, C \rightarrow D, A \rightarrow A, B \rightarrow B, C \rightarrow C, D \rightarrow D, A \rightarrow C, A \rightarrow D, B \rightarrow D\}$$

- Armstrong's Axioms

Reflexivity if $Y \subseteq X, X \rightarrow Y$ (trivial dependency)

Augmentation if $X \rightarrow Y$, then $XZ \rightarrow YZ$

Transitivity if $X \rightarrow Y$ and $Y \rightarrow Z$, then $X \rightarrow Z$

- Additional Rules

Union if $X \rightarrow Y$ and $X \rightarrow Z$, then $X \rightarrow YZ$

Decomposition if $X \rightarrow YZ$, then $X \rightarrow Y$ and $X \rightarrow Z$

- Attribute Closure

~ X^+ - closure of X under F

$$F = \{A \rightarrow B, A \rightarrow C, CG \rightarrow H, CG \rightarrow I, B \rightarrow H\}$$

To compute AG^+ closure = AG $\because A^+ \twoheadrightarrow R$ and $G^+ \twoheadrightarrow R$

closure = ABG $\therefore AG \rightarrow R$

closure = $ABC G$ $\therefore AG$ is a candidate key

closure = $ABC GH$

closure = $ABC GHI$

C. Normal Form

1NF	1st Normal Form	No repeating data group
2NF	2nd Normal Form	No partial key dependency
3NF	3rd Normal Form	No transitive dependency
BCNF	Boyce-Codd Normal Form	Reduce key dependency
4NF	4th Normal Form	No multi-valued dependency
5NF	5th Normal Form	No join dependency

★ $1NF \supseteq 2NF \supseteq 3NF \supseteq BCNF \supseteq 4NF \supseteq 5NF$

- 1st Normal Form

~ each attribute should be only atomic values

- 2nd Normal Form

~ all non-key attributes are functionally dependent on the primary key

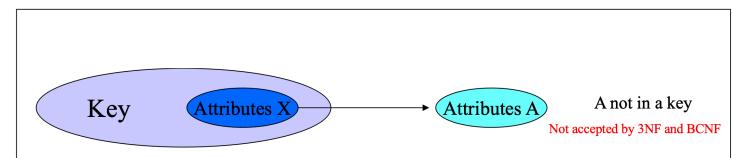
- 3rd Normal Form

~ R is in 3NF if for any $X \rightarrow A$ in F,

1. $X \rightarrow A$ is a trivial dependency ($A \subseteq X$), OR.

2. X is a superkey for R, OR.

3. A is a part of some key for R

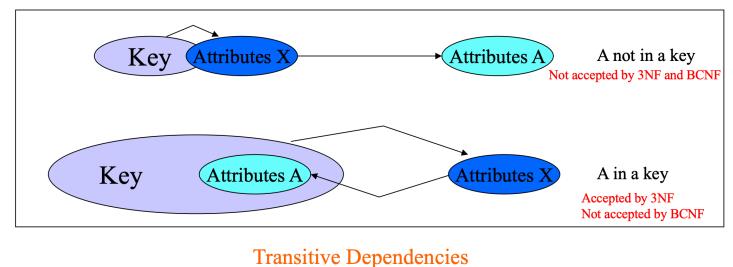


- Boyce-Codd Normal Form

~ R is in BCNF if for any $X \rightarrow A$ in F,

1. $X \rightarrow A$ is a trivial dependency ($A \subseteq X$), OR.

2. X is a superkey for R



D. Decomposition

~ A set of relation schemas $\{R_1, R_2, \dots, R_n\}$ is a decomposition of R if

$$R_1 \cup R_2 \cup \dots \cup R_n = R$$

Supply	<table border="1"> <tr> <td>sid</td><td>status</td><td>city</td><td>part_id</td><td>qty</td></tr> </table>	sid	status	city	part_id	qty
sid	status	city	part_id	qty		

Supplier	<table border="1"> <tr> <td>sid</td><td>status</td><td>city</td></tr> </table>	sid	status	city
sid	status	city		

and

SP	<table border="1"> <tr> <td>sid</td><td>part_id</td><td>qty</td></tr> </table>	sid	part_id	qty
sid	part_id	qty		

$$\Rightarrow Supply = Supplier \cup SP$$

* Decomposition may turn non-normal form into normal form

Suppose R is not BCNF with $FD: X \rightarrow A$ (where $X \cap A = \emptyset \Rightarrow X$ cannot be super key) violates the condition

Step 1: Remove A from R ,

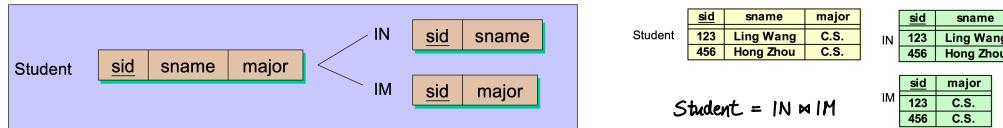
Step 2: Create a new relational schema XA

Repeat until all relations are BCNF

• Lossless Join Decomposition

~ The set of relation schemas $\{R_1, R_2, \dots, R_n\}$ is a lossless join decomposition of R if

for all possible relation instances r on schema R , $r = \Pi_{R_1}(r) \bowtie \Pi_{R_2}(r) \bowtie \dots \bowtie \Pi_{R_n}(r)$



\Rightarrow Student can be recovered by joining the instance of IN and IM

Theorem

the decomposition of R into relations with attribute sets R_1, R_2 is a lossless join decomposition iff

$(R_1 \cap R_2) \rightarrow R_1 \in F^+$ ie. $R_1 \cap R_2$ is a superkey for R_1 or R_2

OR.

the attributes common to R_1 and R_2 must contain a key for either R_1 or R_2

$(R_1 \cap R_2) \rightarrow R_2 \in F^+$

• Dependency Preservation

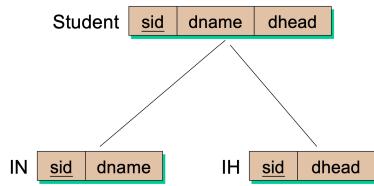
Given relation schema R, Functional Dependencies F and Decomposition $\{R_1, R_2\}$

F_i - the set of dependencies in F^+ involves only attributes in R_i

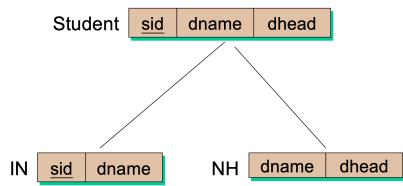
~ the projection of F on the set of attributes in R_i

★ Dependency is preserved if $(F_1 \cup F_2)^+ = F^+$

Dependency set: $F = \{ sid \rightarrow dname, dname \rightarrow dhead \}$



Dependency set: $F = \{ sid \rightarrow dname, dname \rightarrow dhead \}$



$$F_{IN} = \{ \text{trivial dependencies, } sid \rightarrow dname, sid \rightarrow sid, dname \rightarrow dname \}$$

$$F_{IH} = \{ \text{trivial dependencies, } sid \rightarrow dhead, sid \rightarrow sid, dhead \rightarrow dhead \}$$

$$dname \rightarrow dhead \in F^+ \text{ BUT } dname \rightarrow dhead \notin (F_{IN} \cup F_{IH})^+$$

∴ does not preserve dependency

$$F_{IN} = \{ \text{trivial dependencies, } sid \rightarrow dname, sid \rightarrow sid, dname \rightarrow dname \}$$

$$F_{NH} = \{ \text{trivial dependencies, } dname \rightarrow dhead, dname \rightarrow dname, dhead \rightarrow dhead \}$$

$$dname \rightarrow dhead \in F^+ \text{ AND } dname \rightarrow dhead \in (F_{IN} \cup F_{NH})^+$$

∴ preserved dependency

• BCNF Decomposition

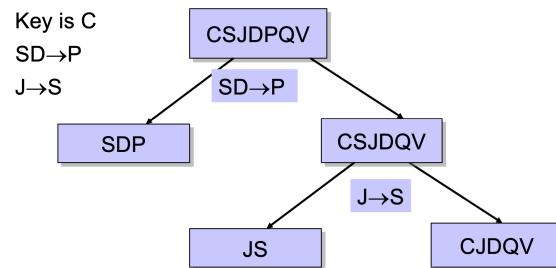
~ Suppose R is not in BCNF, A is an attribute where $X \rightarrow A$ is a FD violates the BCNF condition

Step 1: Decompose R into XA and R-A

Repeat step 1 until all relations are BCNF

★ It is a lessloss join decomposition

★ Not necessary dependency preserving



⇒ Here is a redundancy of CJP

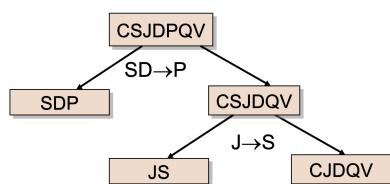
★ Not preserve dependency

Key is C

SD → P

J → S

JP → C



The result is in BCNF

Does not preserve JP → C, we can add a schema: CJP

• Canonical Cover

~ a minimal and equivalent set of functional dependency

equivalent: $E^+ = F^+$ if E is equivalent to F

EG. $R = (A, B, C)$

$$F = \{A \rightarrow BC, B \rightarrow C, A \rightarrow B, AB \rightarrow C\}$$

$\because A \rightarrow BC$ implies $A \rightarrow B$ and $A \rightarrow C$

$\therefore A \rightarrow B$ is redundant

Compute the closure of A , $A^+ = ABC$

$$F' = \{A \rightarrow BC, B \rightarrow C, AB \rightarrow C\}$$

Compute $(AB)^+$ under F' $\because (AB)^+ = ABC$

$\therefore AB \rightarrow C$ is redundant

$$F'' = \{A \rightarrow BC, B \rightarrow C\}$$

\because From $A \rightarrow B$ and $B \rightarrow C$, we deduce $A \rightarrow C$ (transitivity)

\therefore From $A \rightarrow B$ and $A \rightarrow C$, we get $A \rightarrow BC$ (union)

$$F''' = \{A \rightarrow B, B \rightarrow C\}$$

Canonical Cover F_c

1. Every functional dependency $\alpha \rightarrow \beta$ in F_c contains no extraneous attributes in α
2. Every functional dependency $\alpha \rightarrow \beta$ in F_c contains no extraneous attributes in β
3. No two dependencies that $\alpha_1 \rightarrow \beta_1$ and $\alpha_2 \rightarrow \beta_2$ in F_c such that $\alpha_1 = \alpha_2$

Step 1: Replace any $\alpha_1 \rightarrow \beta_1$ and $\alpha_1 \rightarrow \beta_2$ by $\alpha_1 \rightarrow \beta_1\beta_2$

Step 2: Delete any extraneous attributes from any $\alpha \rightarrow \beta$

Repeat until F does not change

- 3NF Synthesis Algorithm

for each $\alpha \rightarrow \beta$ in F_c

if no schema in result contains $\alpha\beta$,

then add schema $\alpha\beta$ to result

if no schema in result contains a candidate key for R,

then choose any candidate key α for R,

add schema α to the result

Ch6. Storage & Indexes

A. Disks and Files

~ DBMS stores information on hard disks

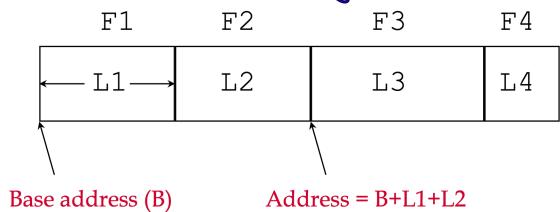
READ: transfer data from disk to main memory RAM

WRITE: transfer data from RAM to disk

~ Disk blocks or pages

page: fixed size and contains a sequence of records

- Record Format: Fixed Length



Ch7. Evaluating Relational Operators

A. Cost Calculations

Example

Schema: Sailors(sid: integer, sname: string, rating: integer, age: real)
Reserves(sid: integer, bid: integer, day: dates, rname: string)

Assumption: Sailors: 500 pages, 80 tuples/page, 50 bytes/tuple
Reserves: 1000 pages, 100 tuples/page, 40 bytes/tuple

Query SELECT *
 FROM Reserves R
 WHERE R.rname = 'Joe'

1. No index, Unsorted (R contains M pages)

1. Scanning the entire R relation by reading the page one by one
2. For each page scanned, checking the condition on each tuple
3. If the condition is met, add the tuple to results

$$\Rightarrow \text{Cost} = M \text{ I/Os}$$

EG. Select rname = 'Joe' from R

$$\Rightarrow \text{Cost} = 1000 \text{ I/Os}$$

2. No index, R sorted on R.attr

1. Binary search on R.attr for value to locate the first tuple that satisfies the condition

$$\sim \text{cost} = \lg M$$

2. Start from that position, scan the relation until the condition become untrue

$\sim \text{cost}$ depends on # tuples satisfying the condition ($10\% \Rightarrow 10\% \cdot M$)

$$\Rightarrow \text{Cost} = \lg M + pM \quad \# p \in [0, 1]$$

EG. Search 'Chan' < rname < 'Lin' and assume 10% of tuples satisfy the condition

$$\Rightarrow \text{Cost} = \lg 1000 + 10\% \cdot 1000 = 110 \text{ I/Os}$$

3. B+ Tree Index on R.attr

1. Search the tree to find the first data entry that points to a qualifying tuple of R

$$\sim \text{cost} = \text{height of the tree} + 1$$

2. Scan the leaf page to retrieve all the data entries in which the key value satisfy the condition

$$\sim \text{cost depends on the number of such data entries}$$

3. For each data entry retrieved, follow the pointer to get the corresponding tuple of R

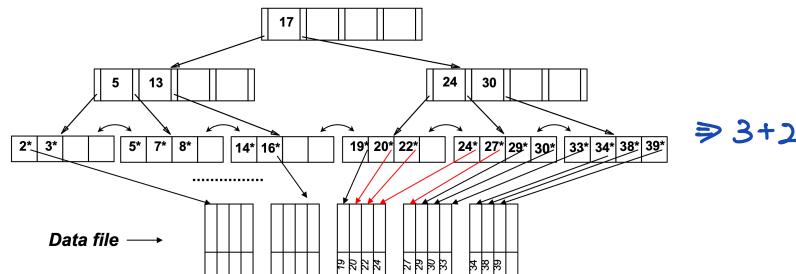
$$\sim \text{let } N \text{ qualifying tuples and } P \text{ tuples that can be stored in a page}$$

$$\text{cost if index is clustered} \leq \lceil N/P \rceil + 1$$

Cluster B+ Tree Index

EG. Search 'Chan' < rname < 'Lin' and assume 10% of tuples satisfy and level = 3

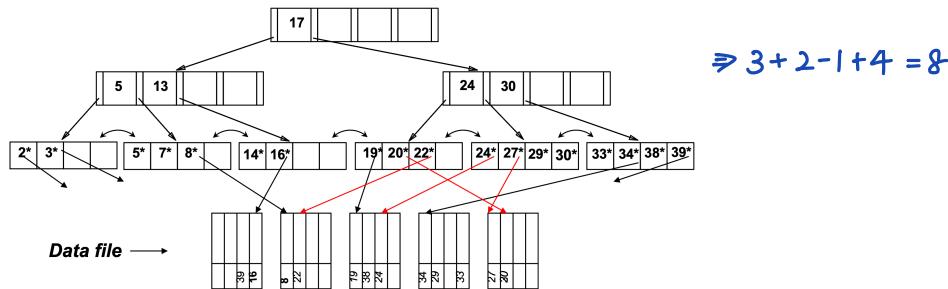
$$\Rightarrow \text{Cost} = 3 + 1000 \times 10\% = 103 \text{ I/Os}$$



Unclustered B+ Tree Index

EG. Assume the total number of pages is $\frac{1}{10}$ of the number of data page

$$\Rightarrow \text{Cost} = 3 + 1000 \times \frac{1}{10} \times \frac{1}{10} = 13 \text{ I/Os}$$



4. Hash index on R.attr (op is =)

1. Calculate the hash value for value

$\sim \text{Cost} = 0$

2. Get the directory entry identified by hash value

$\sim \text{Cost} = 0/1$ | if directory does not fit in the memory (not in main memory)

3. Retrieve the bucket page pointed by the directory entry

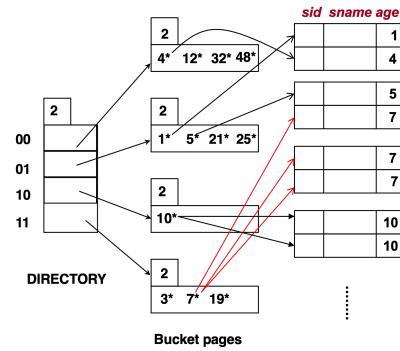
$\sim \text{Cost} = 1/2$ (we may go through the overflow pages)

4. For each data entry in bucket, retrieve the qualifying tuple

$\sim \text{Cost}$ depends on # qualifying tuples and whether clustered index

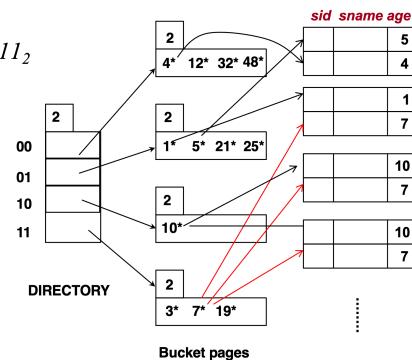
Clustered Index

- Selection Condition: age = 7
- $\mathbf{h}(7) = 7 \bmod 32 = 7 = 111_2$
- Cost
 - Calculate $\mathbf{h}(7)$: 0
 - Get directory entry: 0 (assume directory is in main memory)
 - Get bucket page: 1
 - Get qualifying tuples: 2



Unclustered index

- Sel. Cond. age = 7
- $\mathbf{h}(7) = 7 \bmod 32 = 7 = 111_2$
- Cost
 - Get directory entry: 0 (assume directory is in main memory)
 - Get bucket page: 1
 - Get qualifying tuple: 3



$$2 \times M \times \lceil \log_{B-1} \frac{M}{B} \rceil + 1 \quad M \text{ pages } B \text{ buffers.}$$

B. Projection

~ General Method:

1. Scan relation R and discard the unwanted attributes
2. Eliminate duplicates

~ Projection based on sorting (R - 1000 pages, T - 250 pages, 20 buffer available)

1. Scan R, write the target attributes of each tuple to a temporary file T

$$1000 + 250 = 1250 \text{ I/Os}$$

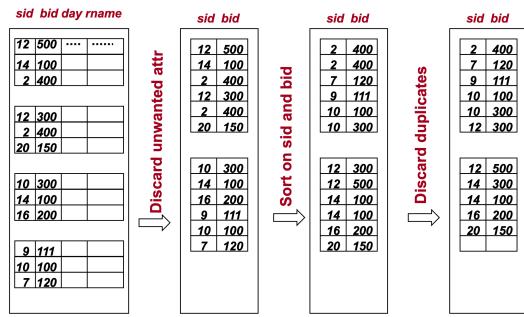
2. Sort T

$$2 \times 250 \times (\lceil \log_{19} \frac{250}{19} \rceil + 1) = 1000 \text{ I/Os}$$

3. Scan the sorted file, compare the adjacent tuples and discard duplicates

$$250 \text{ I/Os}$$

Example: $\Pi_{\text{sid}, \text{bid}}(\text{Reserves})$ based on sorting



C. Join

• Block Nested Loops Join

~ Suppose we have enough buffers to hold $B+2$ pages, then we can read B pages at a time

For each block P of Reserves
for each page q of Sailors
for each $r \in P$ and $s \in q$ such that $r.sid = s.sid$
add $\langle r, s \rangle$ to the result

SELECT *
FROM Reserves R, Sailors S
WHERE R.sid = S.sid

($B=100$)
(R has 1000 pages, 100 tuples/page)
(S has 500 pages, 80 tuples/page)

1. No index

$\therefore B = 100 \Rightarrow$ Reserves can hold $\frac{1000}{100} = 10$ blocks

For each block, scan Sailors for 500 I/Os

\therefore Total Cost = $1000 + 10 \times 500 = 6000$ I/Os

⇒ If we swap them, Total Cost = $500 + \frac{500}{100} \times 1000 = 5500$ I/Os

2. Hash index

Case 1 : Assume an index on sid of the Sailors relation

1. Scan the Reserves relation 1000
 2. For each tuple in Reserves, an average of 1.2 to get the bucket page containing the matching Sailors entry 1.2
 3. Total # Reserves entries 1000×100
 4. # Reserves entries = 100000 while # Sailors = 40000 \Rightarrow each reservations = 0.4 sailors
- \therefore Total Cost = $1000 + 1000 \times 100 \times (1+1.2) = 221000$ I/Os

Case 2 : Clustered Index. Assume an index on sid of the Reserves relation

4. # Reserves entries = 100000 while # Sailors = 40000 \Rightarrow each sailors = 2.5 reservations
- Clustered Index \Rightarrow 2.5 reservations likely on the same page
- \therefore Total Cost = $500 + 500 \times 80 \times (1+1.2) = 88500$ I/Os

Case 3: Unclustered Index

Clustered Index \Rightarrow 2.5 reservations likely NOT on the same page

\therefore Total Cost = $500 + 500 \times 80 \times (2.5+1.2) = 148500$ I/Os

D. Sort - Merge Join

~ Cost for joining two sorted relations are $M + N$

Ch8. Database Concurrency Control

A. Transition

Transition ~ a sequence of read/ write operations a set of transitions $T = \{T_1, T_2, \dots, T_n\}$

History ~ a sequence of operations in T_1, T_2, \dots, T_n $\forall i$

T1	T2
Read[a]	Read[b]
Write[b]	Write[a]

Transaction T1 Transaction T2

1. each operation in T_i appears exactly once
2. operations from T_i appear in the same order as in T_i

Assume: Each T_i is correct when executed individually

T1	T2		
Write[b]			
	Read[b]		
Read[a]			
	Write[a]		

Order of operations is wrong

Invalid Schedule

T1	T2		
Read[a]			
	Write[b]		
		Read[b]	
			Write[a]
	Read[a]		

duplicated operations

Invalid Schedule

T1	T2		
Read[a]			
		Read[b]	
Write[b]			
			Write[a]

Valid schedule (interleaved schedule)

r₁[a] r₂[b] w₁[b] w₂[a]

T1	T2		
Read[a]			
		Read[b]	
Write[b]			
			Read[b]
			Write[a]

Valid schedule (serial schedule)

r₁[a] w₁[b] r₂[b] w₂[a]

• Serial Schedule (Accept Schedule)

~ a schedule is **serial** if $\forall i, j$, either

- all the operations in T_i appear before all the operations in T_j OR
- all the operations in T_j appear before all the operations in T_i

• Equivalent

~ same effect:

the effects of schedule are the value produced by the Write operations of **unaborted transactions**

Assume that if each transaction's Reads read the same value in two schedules,

then all Writes write the same value in both schedules

If each data item x , the final Write on x is the same in both schedules,

then the final values of all data items will be the same in both schedules

B. View Serializability

- View Equivalent

~ Two Schedules H and H' are **view equivalent** if

1. they are over the same set of transactions and have the same operations AND
2. for any un aborted T_i, T_j and for any x , if T_i reads x from T_j in H , then T_i reads x from T_j in H' AND
3. for each x , if $w_i[x]$ is the final write of x in H , then it is also the final write of x in H'

- View Serializable VSR

Assume: there is a transaction T_b that initializes the value for all data objects

~ a schedule is **view serializable** if it is view equivalent to a serial schedule

$r_3[x] w_4[x] w_3[x] w_6[x]$

- T_3 read x from T_b .
- The final write for x is $w_6[x]$.
- View equivalent to $T_3 T_4 T_6$.

$r_3[x] w_4[x] w_3[x]$

- T_3 read x from T_b .
- The final write for x is $w_3[x]$.
- Not serializable.

$r_3[x] w_4[x] r_7[x] w_3[x] w_7[x]$

- T_3 read x from T_b .
- T_7 read x from T_4 .
- The final write for x is $w_7[x]$.
- View equivalent to $T_3 T_4 T_7$.

$w_1[x] r_2[x] w_2[x] r_1[x]$

- T_2 read x from T_1 .
- T_1 read x from T_2 .
- The final write for x is $w_2[x]$.
- Not serializable.

C. Conflict Serializability

~ Non-Conflicting Operations: can always be swapped without affecting the result of the other operations

ie. $op_i(A) op_j(B) \leftrightarrow op_j(B) op_i(A)$ if $A \neq B$

★ Conflict Operations:

1. $r_i[x] r_j[x] \leftrightarrow r_j[x] r_i[x]$

2. $r_i[x] w_j[x]$

3. $w_i[x] r_j[x]$

4. $w_i[x] w_j[x]$

Two operations are conflicting with each other if they both operate on the same data object and at least one of them is a Write

Cannot Swap

• Conflict Equivalent

~ Two histories H and H' are conflict equivalent if

1. they are defined over the same set of transactions and have the same operations
2. they order conflicting operations in the same way

• Conflict Serializable CSR

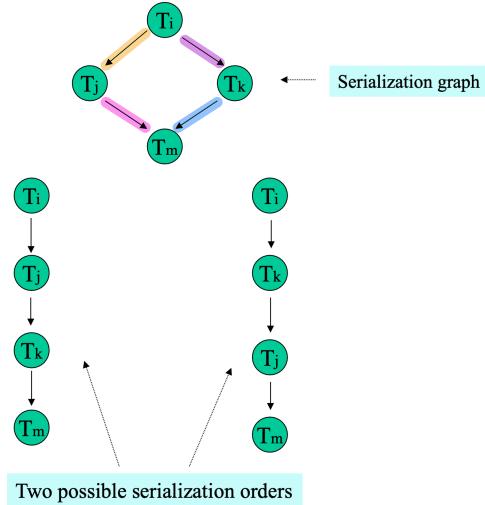
~ a history is conflict serializable if it is conflict equivalent to a serial history

D. Serialization Graph

Recall the Conflict Operations :

- 2. $r_i[x] w_j[x]$
- 3. $w_i[x] r_j[x]$
- 4. $w_i[x] w_j[x]$

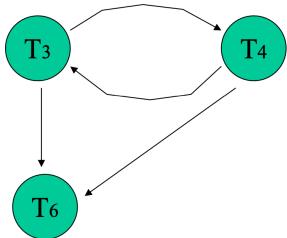
$w_i[a] r_i[b] r_k[a] w_j[b] w_j[d] w_k[c] r_m[c] r_m[d]$



* NOT Conflict Serializable

~ if there is a cycle in the serialization graph

$r_3[x] w_4[x] w_3[x] w_6[x]$



* Conflict serialization \leq View serialization

E. Lock-based Protocols

1. Shared-Locks (read operation)
2. Exclusive-Locks (write operation)

Ch9. Recovery

A. Recoverability

- Recoverable RC

- ~ transitions may be aborted due to logical failure (deadlock)

- ~ recoverability is required to ensure that aborting a transaction does not change the semantics of committed transaction's operations.

EG. $w_1[x] r_2[x] w_2[y] c_2 c_1 \Rightarrow$ Not recoverable: T_2 has committed before T_1

- ~ H is RC if T_i reads from T_j , then $c_j < c_i$

- ie. if each transaction commits after the commitment of all transitions from which it reads

- Avoiding Cascading Aborts

- ~ even for recoverable execution, aborting a transaction may trigger further abortions

- ~ H is ACA if T_i reads x from T_j , $c_j < r_i[x]$

- ie. a transition may read only those values that are written by committed transitions or itself

- Strict Execution

- ~ ACA is not always enough from the practical point of view